


Sub-string/Pattern Matching in Sub-linear Time Using a Sparse Fourier Transform Approach

Krishna R. Narayanan

Department of Electrical and Computer Engineering
Texas A&M University




Problem Statement



- **Database/String:** $\underline{x} = [x[0], x[1], \dots, x[N-1]]$ (length N)
- **Query/Substring:** $\underline{y} = [y[0], y[1], \dots, y[M-1]]$ (length $M = N^\mu$)
- **Signal Model:** $x[i]$'s are i.i.d. r.v. from $\mathcal{A} = \{+1, -1\}$ (extensions possible)

Problem Statement




- **Database/String:** $\underline{x} = [x[0], x[1], \dots, x[N-1]]$ (length N)
- **Query/Substring:** $\underline{y} = [y[0], y[1], \dots, y[M-1]]$ (length $M = N^\mu$)
- **Signal Model:** $x[i]$'s are i.i.d. r.v. from $\mathcal{A} = \{+1, -1\}$ (extensions possible)

Determine all the L locations $\underline{\tau} = [\tau_1, \tau_2, \dots, \tau_L]$ with high probability where

- 1 **Exact Matching:** \underline{y} appears exactly in \underline{x}
 - $\underline{y} := \underline{x}[\tau : \tau + M - 1]$
- 2 **Approximate Matching:** \underline{y} is a noisy substring of \underline{x}
 - $\underline{y} := \underline{x}[\tau : \tau + M - 1] \odot \underline{b}$
 - \underline{b} is a noise sequence with $d_H(\underline{y}, \underline{x}[\tau : \tau + M - 1]) \leq K$

Notations



<i>Symbol</i>	<i>Meaning</i>
N	Size of the string or database in symbols
$M = N^\mu$	Length of the query in symbols
$L = N^\lambda$	Number of matches
K	$\max_{\tau} d_H(\underline{x}[\tau : \tau + M - 1], \underline{y})$
η	$\frac{K}{M}$
$G = N^\gamma$	Number of blocks
$\tilde{N} = N^{1-\gamma}$	Length of one block
$f_i = N^\alpha$	Length of smaller point IDFT at each branch
$g_i = N/f_i$	Sub-sampling parameter
B	Number of shifts also referred to as branches
d	Number of stages in the FFAST algorithm

Main Result

Theorem 1

Assume that a sketch of \underline{x} of size $O(\frac{N}{M} \log N)$ can be precomputed and stored. Then for the exact pattern matching and approximate pattern matching (with $K = \eta M$, $0 \leq \eta \leq 1/6$) problems, with the number of matches L scaling as $O(N^\lambda)$, our algorithm has

- a sketching function for \underline{y} that computes $O(\frac{N}{M} \log N) = O(N^{1-\mu} \log N)$ *samples*
- a *computational complexity* of $O(\max(N^{1-\mu} \log^2 N, N^{\mu+\lambda} \log N))$
- a decoder that recovers all the L matching positions with a *failure probability that approaches zero asymptotically*

Main Result

Theorem 1

Assume that a sketch of \underline{x} of size $O(\frac{N}{M} \log N)$ can be precomputed and stored. Then for the exact pattern matching and approximate pattern matching (with $K = \eta M$, $0 \leq \eta \leq 1/6$) problems, with the number of matches L scaling as $O(N^\lambda)$, our algorithm has

- a sketching function for \underline{y} that computes $O(\frac{N}{M} \log N) = O(N^{1-\mu} \log N)$ *samples*
- a *computational complexity* of $O(\max(N^{1-\mu} \log^2 N, N^{\mu+\lambda} \log N))$
- a decoder that recovers all the L matching positions with a *failure probability that approaches zero asymptotically*

Note

Particularly when $L = O(1)$ or $L < \frac{N}{M}$ (i.e. $\lambda < 1 - \mu$) our algorithm has a *sub-linear* time and space complexity.

Some Prior Work

Exact Matching

- **Boyer and Moore 1977**: First occurrence of the match (only τ_1)
 - Average complexity - $O(N^{1-\mu} \log N)$ (sublinear)
 - Worst case complexity - $O(N \log N)$
- **Goodrich, Atallah, and Tamassia 2005**: BWT, suffix-arrays based indexing
 - Time complexity - $O(M + L)$ (sublinear)
 - Storage Complexity - $O(N H_k(X) \log^\epsilon N) + o(N)$ bits (linear)
 - Read alignment in Bio-informatics community[[li2009fast](#); [li2010fast](#)]

Approximate Matching

- **Chang and Marr 1994**: Generalization of Boyer and Moore 1977
 - Average time complexity - $O(NK/M \log N)$ (sub-linear only when $K \ll M$)
- **zhang2003approximate**: Approximate Matching using BWT
 - Worst case time complexity: $O(\min\{M(M-K)|\mathcal{A}|^k \log \frac{N}{|\mathcal{A}|}, NM \log \frac{N}{|\mathcal{A}|}\})$
 - Complexity grows with $|\mathcal{A}|$ and K
- **Andoni et al. 2013**: $O(N/M^{0.359})$ (sub-linear even when $K = O(M)$)
 - Combinatorial in nature

Some Prior Work

Sparse Fourier Transform Approach

- [Pawar and Ramchandran 2014](#): Robust Sparse Fourier Transform
 - Sparse graph code approach
 - Computational complexity : $O(N \log N)$
- [Hassanieh et al. 2012](#): Faster GPS receiver
 - Exploited sparsity in Correlation function R_{XY}

Motivation

- **Cross-correlation** (\underline{r}):

$$r[m] = (\underline{x} * \underline{y})[m] \triangleq \sum_{i=0}^{M-1} x[m+i]y[i], \quad 0 \leq m \leq N-1$$


- **Naive implementation:** $O(MN) = O(N^{1+\mu})$ (**super-linear** complexity)
- **Fourier Transform Approach:** $O(N \log N)$ complexity

$$\underline{r} = \mathcal{F}_N^{-1} \{ \mathcal{F}_N \{ \underline{x} \} \odot \mathcal{F}_N \{ \underline{y}' \} \}, \quad \underline{y}' = \underline{y}^*[-n]$$

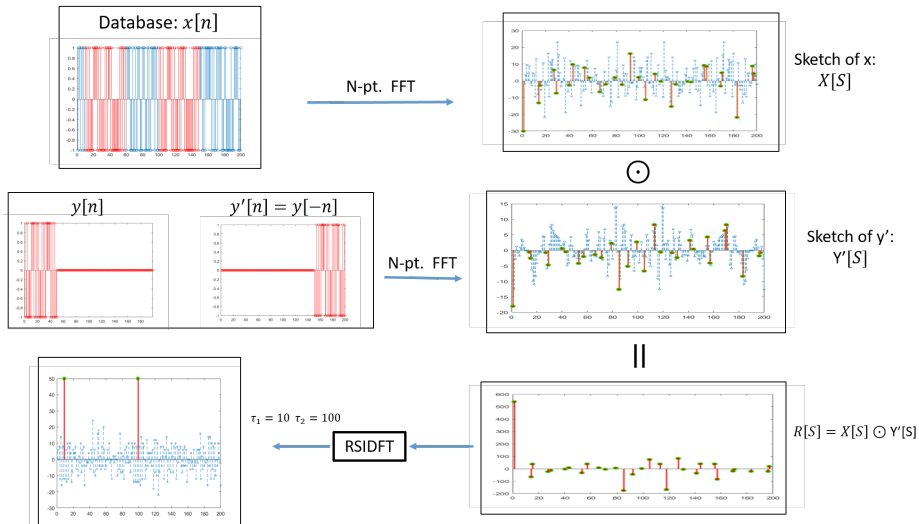
Key Observation

- \underline{r} is **Sparse** with some noise.


$$r[m] = \begin{cases} M, & \text{if } m \in \mathcal{T} \\ n_m, & m \in [N] - \mathcal{T} \end{cases}$$



Example



Sparse Fourier Transform Approach



$$\underline{r} = \underset{3}{\mathcal{F}_N^{-1}} \left\{ \underset{1}{\mathcal{F}_N\{\underline{x}\}} \odot \underset{2}{\mathcal{F}_N\{\underline{y}'\}} \right\}$$

1. **Sketch of \underline{x}** : Assume $\underline{X}[l] = \mathcal{F}\{\underline{x}\}$ is precomputed at positions $l \in S$.
2. **Sketch of \underline{y}** : Compute $\underline{Y}'[l] = \mathcal{F}\{\underline{y}'\}$ for $l \in S$.
 - Only M non-zero values in \underline{y}' - Efficient computation (folding and adding)
3. **Sparse \mathcal{F}^{-1}** :
 - Robust Sparse Inverse Fourier Transform (RSIDFT)
 - Efficient Implementation- **sublinear** time and sampling complexity

Robust Sparse Inverse Fourier Transform(RSIDFT)

Main Idea


- **Sub-sampling** in frequency corresponds to **aliasing** in time
- Aliased coefficients \Leftrightarrow parity check constraints of **GLDPC codes**
- **CRT** guided sub-sampling induces a code good for **Peeling decoder**
- R-FFAST- proposed by Pawar and Ramchandran 2014

Key modifications


- Optimized for the induced noise model
- Correlation peak is always **positive**
- Take advantage in decoding algorithm - **sub-linear** time complexity

Aliasing and Sparse Graph Codes


IDFT Computation ($N = 6$)




x_s : Sub-sampled by $P_1 = 2$



z_s : Sub-sampled by $P_2 = 3$




Factor graph




Aliasing and Sparse Graph Codes


IDFT Computation ($N = 6$)




\tilde{x}_s : Sub-sampled by $P_1 = 2$ (shifted)




\tilde{z}_s : Sub-sampled by $P_2 = 3$ (shifted)





Factor graph




RSIDFT Framework



RSIDFT Framework



RSIDFT-Decoding (Peeling Decoder)



$$w_i^k := \left(e^{j \frac{2\pi s_i}{N}} \right)^k$$

Observations:

$$z_{i,k} = \begin{bmatrix} r_{i,1}[k] \\ r_{i,2}[k] \\ \vdots \\ r_{i,B}[k] \end{bmatrix}^T$$

Decoding- 3 steps

- 1 Bin Classification
- 2 Position Identification
- 3 Peeling Process

Decoder

Bin Classification

- Classify each check-node - Zero-ton / Single-ton / Multi-ton
- **Threshold constraints** on first observation $z_{i,k}[1] = z$
- Threshold varies with η
 - different for exact($\eta = 0$) and approximate matching

$$\hat{\mathcal{H}}_{i,j} = \begin{cases} \mathcal{H}_z & z/M < \gamma_1 \\ \mathcal{H}_s & \gamma_1 < z/M < \gamma_2 \\ \mathcal{H}_d & \gamma_2 < z/M < \gamma_3 \\ \mathcal{H}_m & z/M > \gamma_3 \end{cases}$$

where $(\gamma_1, \gamma_2, \gamma_3) = (\frac{1-2\eta}{2}, \frac{3-4\eta}{2}, \frac{5-6\eta}{2})$

Decoder

Position Identification

- Observation:

$$\underline{z}_{i,k} = \begin{bmatrix} 1 & 1 & \dots & 1 \\ \omega^{ks_2} & \omega^{(k+f_i)s_2} & \dots & \omega^{(k+(g_i-1)f_i)s_2} \\ \vdots & \vdots & \ddots & \vdots \\ \omega^{ks_B} & \omega^{(k+f_i)s_B} & \dots & \omega^{(k+(g_i-1)f_i)s_B} \end{bmatrix} \times \begin{bmatrix} r[k + (0)f_i] \\ r[k + (1)f_i] \\ \vdots \\ r[k + (g_i - 1)f_i] \end{bmatrix}$$

- Column that gives **maximum correlation** with the observation

$$\hat{k} = \arg \max_{k \in \{j+lf_i\}} \underline{z}_{i,j}^\dagger \mathbf{W}[:, l]$$

Decoder


Peeling Process:

Exact Matching

- Remove a decoded variable node's contribution from **all participating bin nodes**

Approximate Matching

- Remove a decoded variable node's contribution only from neighboring **single-tons and double-tons**
- Avoid error propagation




Error Analysis

Error Events

- \mathcal{E}_1 -Bin Classification: Bin is wrongly classified
- \mathcal{E}_2 -Pos. Identification: Position of singleton is identified wrongly, given a singleton
- \mathcal{E}_3 -Peeling Process: Peeling process fails to recover the L significant correlation coefficients, given $\mathbb{P}(\mathcal{E}_1) = \mathbb{P}(\mathcal{E}_2) = 0$

\mathcal{E}_1 -Bin Classification



$$\begin{aligned}
 \mathbb{P}[\mathcal{E}_1] &\leq \mathbb{P}[\mathcal{E}_1 | \hat{\mathcal{H}}_{i,j} = \mathcal{H}_z] + \mathbb{P}[\mathcal{E}_1 | \hat{\mathcal{H}}_{i,j} = \mathcal{H}_s] + \mathbb{P}[\mathcal{E}_1 | \hat{\mathcal{H}}_{i,j} = \mathcal{H}_d \cup \mathcal{H}_m] \\
 &= \mathbb{P}[z[1] > \gamma_1] + (1 - \mathbb{P}[\gamma_1 < z[1] < \gamma_2]) + \mathbb{P}[z[1] < \gamma_2]
 \end{aligned}$$

Error Analysis

\mathcal{E}_2 -Pos. Identification

- $\underline{z} = r[j_p] \underline{w}_{j_p} + \sum_{k \neq p} n_k \underline{w}_{j_k}$
- $\mathbb{P}[\mathcal{E}_2] = \mathbb{P}[\underline{w}_{j_p}^\dagger \underline{z} < \underline{w}_{j_k}^\dagger \underline{z}]$
- Mutual Incoherence property to bound the cross-correlation(noise) term
 - $\log N$ measurements (shifts) suffices [PR2014]

\mathcal{E}_3 -Peeling Process

- Tools from Coding Theory to analyze Sparse Graph Codes
- Density Evolution to quantify Error Probability
- # of check-nodes is a function of sparsity (query length)
- Exponentially decaying error probability- R-FFAST and SAFFRON [PR2014,LPR2015]

Error Analysis

Error Events

- \mathcal{E}_1 -*Bin Classification*: Bin is wrongly classified
- \mathcal{E}_2 -*Pos. Identification*: Position of singleton is identified wrongly, given a singleton
- \mathcal{E}_3 -*Peeling Process*: Peeling process fails to recover the L significant correlation coefficients, given $\mathbb{P}(\mathcal{E}_1) = \mathbb{P}(\mathcal{E}_2) = 0$

Error Probability

$$\begin{aligned}\mathbb{P}(\mathcal{E}_{\text{total}}) &\leq \mathbb{P}(\mathcal{E}_1) + \mathbb{P}(\mathcal{E}_2) + \mathbb{P}(\mathcal{E}_3) \\ &\leq 6e^{-\frac{N\mu+\alpha-1(1-6\eta)^2}{16}} + 2e^{-N\mu+\alpha-1} c_1(\eta) + e^{-c_3 N^{c_4 \alpha}}\end{aligned}$$

$$\boxed{\mathbb{P}(\mathcal{E}_{\text{total}}) \rightarrow 0 \text{ if } \alpha > 1 - \mu}$$

Complexity Analysis

Sample Complexity

$$\text{Total \# of samples required (S)} = O(dBN^\alpha) = O(N^{1-\mu} \log N)$$

Computational Complexity

$$\underline{r} = \mathcal{F}_N^{-1} \{ \underbrace{\mathcal{F}_N\{\underline{x}\}}_{II} \odot \underbrace{\mathcal{F}_N\{\underline{y}'\}}_I \}$$

- Sketch of Query:

$$C_I = dB \left(\underbrace{N^\mu}_{\text{Folding}} + \underbrace{N^\alpha \log N^\alpha}_{\text{Shorter FFTs}} \right) = O(\max(N^{1-\mu} \log^2 N, N^\mu \log N))$$

- RSIDFT:

$$C_{II} = dB \left(\underbrace{O(N^\alpha \log N^\alpha)}_{\text{Shorter IFFT's /block/stage}} + \underbrace{L N^{1-\alpha}}_{\text{Correlations}} \right) = O(\max(N^{1-\mu} \log^2 N, N^{\mu+\lambda} \log N))$$

$$C_{\text{total}} = \max(C_I, C_{II}) = O(\max(N^{1-\mu} \log^2 N, N^{\mu+\lambda} \log N))$$

Simulation Results




Figure: Plot of Probability of Missing a Match vs. Sample Gain for Exact Matching of a substring of length $M = 10^5$ ($\mu = 0.41$) from a equiprobable binary $\{+1, -1\}$ sequence of length $N = 10^{12}$, divided into $G = 10^5$ blocks each of length $\tilde{N} = 10^7$. The substring was simulated to repeat in $L = 10^6$ ($\lambda = 0.5$) locations uniformly at random.

Simulation Results




Figure: Plot of Probability of Missing a Match vs. Sample Gain for Exact Matching of a substring of length $M = 10^3$ ($\mu = 0.25$) from a equiprobable binary $\{+1, -1\}$ sequence of length $N = 10^{12}$, divided into $G = 10^6$ blocks each of length $\tilde{N} = 10^6$. The substring was simulated to repeat in $L = 10^6$ ($\lambda = 0.5$) locations uniformly at random.

Questions?



Thank you!